

WHITE PAPER

Governing Meaning at Scale

The Holonic Semantic Governance Architecture — A Reference Model for Trustworthy Enterprise Knowledge and AI Governance

Based on — Holonic Semantic Governance Architecture v4.2, Reference Topology & Governance Diagram Edition

Author — Jeffrey Wallk, Senior Enterprise Architect & Knowledge Engineering Leader

Organization — The Value Enablement Group, LLC

*Prepared for enterprise and technology leadership
August 2026*

Executive Summary

Every enterprise now runs two systems at once: the one that stores and moves data, and the one — usually informal, usually inconsistent — that decides what that data **means**, who may act on it, and under what conditions. As long as people made most of the decisions, that informality was tolerable. It is not tolerable once AI agents are making decisions, executing processes, and generating content at machine speed, because an AI system cannot infer governance intent that was never made explicit. It can only follow — or violate — whatever structure it is actually given.

The **Holonic Semantic Governance Architecture (HSGA)** is a reference architecture for making that structure explicit. It governs what enterprise data and knowledge mean, what they may be used for, and what actions they may authorize — at design time and at run time, for humans and AI agents alike. It is built on a single organizing idea borrowed from systems theory: Arthur Koestler's concept of the **holon** — an entity that is simultaneously a self-contained whole and a governed part of something larger. HSGA applies that idea consistently, from the smallest governed fact in the enterprise up through the full federated governance ecosystem, and in doing so unifies nine previously separate governance, AI-trust, and systems-engineering frameworks into one coherent architecture.

Why this matters to leadership — *organizations that scale AI without scaling governance discover the cost later, usually as a compliance failure, a hallucinated decision, or a re-engineering bill when a data platform changes underneath a brittle integration. HSGA is built specifically against those three failure modes: it makes compliance a structural property rather than a policy document, it contains AI failure at the level where it occurs instead of letting it cascade, and it insulates the enterprise's meaning from the churn of its underlying technology stack.*

What this paper covers. This paper explains HSGA's four governance levels and the lateral Mediation Holon that connects them; the mechanisms that make the architecture trustworthy for AI use — the Circuit Breaker trust monitor, compiled governance for production-speed enforcement, and a single human-in-the-loop review cycle that runs through every layer; the open standards it is built on; and the strategic value it creates once knowledge and AI activity are governed this way rather than merely logged.

Part I — Why Conventional Enterprise Architecture Breaks at AI Scale

Most enterprise data architecture is organized around a question that sounds reasonable and turns out to be unanswerable at scale: *which system is the single source of truth?* In a large enterprise, no such system exists. Different domains legitimately need authority over the same concept in different contexts — finance owns the authoritative revenue figure for reporting, sales owns it for forecasting, and both are correct within their own governance scope. Architectures built around a global Authoritative Source of Truth (ASOT) do not resolve this; they paper over it, and the seams show up later as data-quality incidents, conflicting dashboards, or compliance gaps nobody can explain.

A second, related failure is **global schema brittleness**. When governance rules are embedded in one monolithic enterprise schema, a legitimate local change — a new product line, a new regulatory category, a new data source — risks breaking the whole schema, so it either gets forced through a slow central review process or it gets built as a workaround outside the governed system entirely. Either outcome erodes the thing the architecture was supposed to protect.

A third failure is specific to model-based systems engineering: **model staleness**. Systems and process models built at design time diverge from operational reality the moment the system starts changing, and without a structural mechanism to keep them synchronized, the model becomes a documentation artifact rather than something that actually governs operations.

The fourth failure is the one that makes the first three urgent rather than merely inconvenient: **ungoverned AI compounds all of them**. An AI system operating without an explicit governance structure produces a well-documented failure sequence — noise or quality problems at the data layer get transformed by inflexible algorithms into false features, an LLM finds plausible-sounding meaning in that erroneous data and hallucinates confidently, and if that hallucinated output is not caught before it re-enters training data, the error compounds recursively into model collapse. Each stage amplifies the one before it, and without a governance layer built to contain failure at the level where it occurs, a local data-quality problem can end several steps later as an autonomous agent taking a materially wrong action.

The reframe HSGA is built around — *the goal is not a faster or better-labeled version of conventional data management. It is a categorically different objective, one Kurt Cagle calls Digital Orchestration rather than Digital Documentation: a system that does not simply record what happened, but governs what should happen and ensures that what actually happens conforms to the governance rules established for it, continuously and at every level of the enterprise.*

Part II — The Architectural Grammar: What a Holon Is

Arthur Koestler introduced the term *holon* in *The Ghost in the Machine* (1967) to name a property he observed across complex natural and social systems: every stable unit of organization is simultaneously a self-contained whole and a dependent part. Holons are neither purely atomic nor purely monolithic. They form nested structures — holarchies — in which local self-governance and global interoperability coexist through shared protocols rather than centralized control.

Three properties define a holon, and HSGA requires all three of any entity it governs:

- **Autonomy** — the holon governs itself according to its own internal invariants without requiring continuous external instruction, and can continue operating in a degraded, “island mode” if higher-level connectivity is unavailable.
- **Integration** — the holon exposes standardized interfaces that let it compose with other holons through governance protocols, never through shared or coupled schemas.
- **Identity** — the holon maintains a stable self-description across contexts, so its meaning and governance rules travel with it wherever it is composed into a larger structure.

Kurt Cagle's concept of *holonic graphs* extends Koestler's theory directly into knowledge representation, and names the failure mode described in Part I precisely: global schema brittleness. His extension contributes three operating principles that recur throughout this architecture — **resilience through semi-autonomy** (distributed governance across component holons means a change in one niche domain never forces a global re-engineering effort), **comprehension over processing** (an agent traversing a holonic graph is navigating governed meaning, not retrieving inert data — context and governance are inseparable from the content itself), and **Digital Orchestration versus Digital Documentation**, introduced in Part I.

This is the single grammar that HSGA applies everywhere. The component-ontology pattern, the semantic gateway's staged traversal, the federated governance hierarchy, the construction/integration/production lifecycle, the trust-monitoring structure, and the AI interaction protocol described in the parts that follow are not nine separate design decisions — they are one architectural principle, distributed self-governance through protocol-coordinated composition, expressed at nine different points in the enterprise. HSGA's own contribution is to unify nine previously independent frameworks — spanning semantic gateway design, AI trust monitoring, federated governance, systems-engineering knowledge standards, and knowledge-graph learning — under that single grammar, so that each pattern reinforces the others instead of operating as an unrelated add-on.

Part III — The Architecture at a Glance

HSGA organizes governance across four vertical levels, plus one lateral element — the Mediation Holon — introduced in Version 3.0. The levels are not organizational reporting tiers; they are structural boundaries at which different governance mechanisms apply, and each level is simultaneously self-governing (a whole) and a governed component of the level above it (a part).

Level	Structural Scope	Governance Mechanism	Key Artifacts
L1 — Knowledge Holon	An individual semantic asset: a vocabulary term, taxonomy node, rule, or constraint	SHACL constraint, SWRL rule, ORSD definition, KGCL change record	OWL class, RDF triple, SHACL shape, SKOS concept
L2 — Domain Holon	A component ontology plus its associated AI agent, vector database, and Policy Enforcement Points	Domain governance team, data contracts, domain-specific SHACL/SWRL	Domain ontology, agent registry, vector index, PEP rule set
L3 — Gateway Holon	The semantic gateway — the holarchic traversal and composition protocol	STPA/STAMP safety coverage, Requirement Context Bundle, Compiled AI Directive	Rule Harmonization output, PEP enforcement engine, Circuit Breaker
L4 — Enterprise Holarchy	The full federated governance ecosystem across every domain and lifecycle stage	Enterprise Governance Council, Center of Excellence, continuous engineering loop	Operating model, DEV/STG/PRD environments, governance observatory
Mediation Holon (lateral)	The correspondence between top-down meaning and bottom-up implementation, not the knowledge itself	Governed mapping graph, Rule Harmonization Engine, Rule Projection Agents	Concept mappings, Enterprise Entity IDs, compiled runtime policy

Figure 1 shows how these levels sit within a complete solution topology — from the physical and operational reality at the bottom of the enterprise, through the knowledge and governance layers this architecture is chiefly about, up to the applications and human consumers at the top. Four governance domains — Data, Semantic, Analytics, and AI/Trust — run as an ambient band across every layer, enforced at each layer boundary by the Gateway Holon rather than applied as an afterthought at the edge.

Full Solution Topology (1 of 2)

The upper holarchy — knowledge, governance, and enterprise applications

L7 — Enterprise Ecosystem

- ERP / financial systems of record
- BI analytics & intelligence applications
- Quality & compliance command center

L6 — Holon Orchestration & UI

- Workflow & orchestration platform
- Ontology & context explorers
- Role-based dashboards

L5 — Semantic Governance

- OWL / RDF domain ontology
- SHACL / SWRL constraints & rules
- Deontic norm library (LKIF)

L4 — Knowledge Graph + Data Fabric

- Property graph, named graphs
- Federated SPARQL / GraphQL / Cypher
- Vector / RAG embedding search

Continues below with L3–L1, the enforcement, connectivity, and physical layers. Four-domain governance (Data · Semantic · Analytics · AI/Trust) is ambient across every layer.

Figure 1a: Full Solution Topology (1 of 2) — the upper holarchy: knowledge, governance, and enterprise applications (L7–L4).

Full Solution Topology (2 of 2)

The lower holarchy — enforcement, connectivity, and physical reality

L3 — PEP & AI Edge — Gateway Holon

- Six-stage requirement-driven traversal
- Compliance & safety enforcement
- Compiled AI Directive assembly

L2 — Digital Twin & Connectivity

- IoT / OPC-UA / API adapters
- Batch ETL & ingestion pipelines
- Event streaming (topic-per-domain)

L1 — Physical / Operational Holons

- Equipment, machines, physical assets
- Sensors, instruments, telemetry
- Source / operational systems of record

Continued from above (L7–L4). Governance is enforced at every layer boundary by the Gateway Holon (L3).

Figure 1b: Full Solution Topology (2 of 2) — the lower holarchy: enforcement, connectivity, and physical reality (L3–L1).

A note on the two level schemes — *the four L1–L4 levels above are the governance-authority axis: they answer who has the right to govern a given piece of meaning. Figure 1's seven layers are a complementary solution-topology view: they answer where governance actually executes, physically and technically, from sensors and source systems up through the applications people and other systems consume. The Gateway Holon (L3 in the authority axis) is the enforcement point common to both views — every layer boundary in Figure 1 is a governance checkpoint it owns, not a data hop.*

Part IV — Levels 1 and 2: Knowledge Holons and Domain Holons

IV.1 Level 1 — every semantic asset is a governed unit

At the finest grain of the architecture, every semantic asset — a vocabulary term, a taxonomy node, a SHACL constraint, an OWL class definition, a SWRL rule, a data-contract clause — is a knowledge holon. It carries its own identity (a persistent URI/IRI), its own governance attributes (version, lifecycle state, steward, provenance), and its own invariants: the formal rules that define what a valid instance of that concept looks like.

Two protocols together govern how a knowledge holon changes over time without losing its identity. The **Ontology Requirements Specification Document (ORSĐ)** is the design-time specification of a holon's purpose, scope, and governance requirements — the holonic equivalent of a requirements baseline in systems engineering. The **Knowledge Graph Change Language (KGCL)** is the run-time change record: every add, update, or deletion is expressed as a KGCL record carrying what changed, why (business rationale and semantic context), who made the change and when, and how any conflict was resolved. A knowledge holon can evolve — its definition can change, its relationships can extend, its constraints can be revised — without ever losing its identity or breaking its governance contract with the rest of the architecture.

The most consequential extension at this level is **LKIF deontic integration**: mapping knowledge holons to Legal Knowledge Interchange Format statements, which elevates them from descriptive units to normative ones — artifacts that do not merely describe what something is, but prescribe what must, may, and must not happen.

LKIF Statement	Holonic Governance Role	Worked Example (Financial Planning & Analysis)
Obligation	A holonic invariant — a rule that must be satisfied for the holon to be in a valid state	Forecast variance must remain below 5% for SOX compliance
Permission	A holonic boundary condition — the envelope within which the holon may vary without intervention	Forecast adjustments within a 2% range are permitted without escalation
Prohibition	A holonic integrity constraint — an action that, if taken, invalidates the holon's governance status	Misstating a forecast is prohibited outright
Violation	The holonic fault state reached when an obligation or prohibition is breached	Triggers Circuit Breaker evaluation at the domain holonic level
Sanction	The holarchic response protocol once a violation is detected	Circuit Breaker trip, mandatory human authorization, or a KGCL change record

This pattern lets agentic AI reason directly over compliance constraints as first-class properties of the data it is working with, rather than as an external policy document the agent has no structural relationship to.

IV.2 Level 2 — component ontologies as sovereign governance units

At Level 2, the component ontology is the fundamental unit. Each domain holon is a self-contained governance environment: a domain ontology (vocabularies, taxonomies, classes, constraints, rules), a domain AI agent with its own vector database indexed by that ontology, and a set of Policy Enforcement Points governing how the domain holon interacts with the rest of the architecture across six domains — Compliance, Risk, Legal, Security, Data Governance, and AI Governance.

Domain specificity is treated as a feature of the architecture, not a constraint to be minimized. In conventional monolithic ontology designs, adding domain-specific concepts risks cascading constraint violations elsewhere; in the holonic design, domain specificity is bounded within the domain holon's own sovereignty, and changes propagate through governance protocol rather than through shared-schema coupling. This is the INCOSE Knowledge Systems Working Group's "Context is King" principle applied structurally: core vocabularies and taxonomies form the enterprise-level protocol every domain must conform to, and each engineering or business discipline then owns its own vocabulary, taxonomy, and ontology stack beneath that — inheriting from the core holon polymorphically, never by schema coupling.

A construction/integration/production lifecycle (DEV/STG/PRD) governs how domain holons mature: DEV is where domain holons are built and tested in isolation; STG is where top-down and bottom-up artifacts are reconciled; PRD is where composed domain holons are published for downstream consumption. One five-dimensional framework — WHO, WHAT, WHEN, WHERE, WHY — recurs at every one of these stages, doing three jobs at once: it is the context vector for every governed interaction, the dimensional axis along which the Circuit Breaker trust metric (Part VI) is stratified, and the governance axes of the semantic gateway's Requirement Context Bundle (also Part VI). A single failure in, say, the WHERE dimension (a domain-routing problem) is diagnosable and repairable without ever touching the WHAT dimension (concept resolution) in an unrelated domain — targeted remediation instead of systemic shutdown.


Part V — The Mediation Holon: Decoupling Meaning from Implementation

Every enterprise semantic layer has to reconcile two things that change at completely different rates. The top-down side — the enterprise ontology, its core vocabularies and domain classes — evolves slowly and deliberately, under governance-council authority. The bottom-up side — property graphs, transformation pipelines, vector databases used for retrieval-augmented generation, third-party ontologies, data catalogs — evolves quickly and constantly, under engineering-team authority. Point-to-point integration between the two is the brittle pattern that has historically absorbed the most enterprise integration effort, because every bottom-up change requires a corresponding top-down fix.

The Mediation Holon, introduced in Version 3.0, resolves this by insulation rather than integration. It is a persistent, governed Mapping Graph — itself defined with the same identity, invariants, and lifecycle protocol as any other holon in the architecture — that sits between the two tiers and absorbs the bottom-up churn so the top-down ontology never has to.

The Three-Tier Semantic Architecture as Holonic Mediation

Insulation, not integration — the Mediation Holon absorbs bottom-up churn so the top-down ontology never has to



Physical changes in any bottom-up system require only a mapping-graph update. The core enterprise ontology, and everything built on it, remains untouched.

Figure 2: The Three-Tier Semantic Architecture as Holonic Mediation.

The practical payoff is direct: a physical change in any bottom-up system — a dbt model gets refactored, a vector database gets re-indexed, a vendor API changes shape — requires only an update to the mapping graph. The core enterprise ontology, and everything built on top of it (semantic-gateway traversal, Policy Enforcement Points, trust scoring), remains untouched. The Mediation Holon is also what makes retrieval-augmented generation trustworthy rather than merely fluent: a vector database alone retrieves text that is **semantically close** to a query with no structural guarantee it refers to the correct real-world entity, or that surfacing it is even permitted. The Mediation Holon closes that gap in three steps — an embedding retrieves the relevant chunk, the Mapping Graph anchors that chunk to a canonical enterprise entity via its Identity & URI Resolution Registry, and only then do deterministic governance rules (not the language model) decide whether the anchored content is permitted to surface to this requester, in this context, for this purpose. Grounding happens before governance evaluation, and governance evaluation happens before content ever reaches the AI agent.


Compiled Semantic Governance — enforcing rules without the performance cost

An instinct once rules live in a formal ontology is to evaluate them there directly, in place, on every request. That does not scale: full SHACL validation or SWRL/OWL reasoning on every read, across every query modality an enterprise actually uses (SQL, GraphQL, SPARQL, vector search), is frequently an NP-hard problem, and forcing a graph engine to compute it on every request produces exactly the performance bottleneck and vendor lock-in a mediation layer is meant to avoid. Removing the formal rules entirely is not an option either — SHACL shapes do not just validate data quality; they define entity boundaries and class structure that a purely procedural policy engine cannot reproduce, and symbolic SWRL logic gives AI reasoning a deterministic, auditable ground truth that a black-box policy engine cannot certify.

The resolution is a clean separation of rule *definition* from rule *execution* — Compiled Semantic Governance. Canonical SHACL shapes, SWRL rules, and deontic norms remain the architecture's single source of truth, authored and version-controlled by domain governance teams. A Rule Harmonization Engine runs conflict and dependency detection across that canonical rule set before anything is compiled. A set of Rule Projection Agents then compile the surviving rules into runtime-native enforcement targets — SQL/dbt constraints, OPA/Rego policy bundles, vector-filter configurations, and agent guardrail schemas — so that production enforcement runs at microsecond latency against compiled, pre-validated targets, never against live graph reasoning.

SHACL/SWRL as Single Source of Truth

The hybrid compiled-governance pattern



Every compiled policy artifact carries a URI back to the canonical SHACL/SWRL invariant that produced it — full audit traceability, no runtime graph reasoning.

Figure 3: SHACL/SWRL as Single Source of Truth — the hybrid compiled-governance pattern that separates definition-time authoring from runtime-native enforcement.

The audit-traceability guarantee — every compiled policy artifact — a dbt test, a Rego bundle, a vector filter, an agent guardrail — carries an explicit link back to the canonical SHACL or SWRL invariant that produced it. A compliance reviewer can always answer “which governance rule authorized this behavior?” without inspecting runtime code.

Part VI — Level 3: The Gateway Holon and the Circuit Breaker

VI.1 A protocol, not a pipeline

The most consequential reframing holonic theory brings to the semantic gateway is the distinction between a pipeline and a holarchic protocol. A pipeline is a sequential transformation: input enters one end, output exits the other, and each stage transforms the content. A holarchic protocol is a governance composition mechanism: context enters, is enriched by the governance rules of every holonic level it traverses, and exits as a governed behavioral directive carrying the semantic authority of every domain holon it touched.

Reframed this way, the gateway's traversal runs in six stages. A **Requirement Context Bundle** is assembled at entry, decomposing the request into its five relevancy dimensions (WHO/WHAT/WHEN/WHERE/WHY) and activating the relevant domain holons. Four enforcement layers then evaluate in sequence — SHACL constraint validation, SWRL temporal-rule application, Policy Enforcement Point compliance evaluation across all six governance domains, and Unsafe Control Action coverage verification against the architecture's STPA/STAMP-based safety holon — each now evaluating pre-compiled rule projections from the Mediation Holon rather than performing full graph reasoning inline. Governance outputs from all four layers are synthesized into a **Compiled AI Directive**: the agent does not receive instructions; it receives governed context. Finally, at the Service Facade, the agent's response is validated against acceptance criteria before it leaves the governance boundary, and irreversible actions are suspended at a Simplex Action Gateway pending explicit human authorization.

Why a protocol cannot be bypassed the way a filter can — *a content filter operates on top of a system and can be defeated through prompt injection or context manipulation, because the filter is a layer, not the structure itself. A holarchic protocol cannot be bypassed the same way, because governance is the structure through which context has to compose to become an actionable directive at all. This is the architecture's operational definition of trustworthy AI at enterprise scale.*

VI.2 Containing failure at the level where it occurs

Part I described how an ungoverned AI system's failures cascade — data-layer noise becomes an algorithmic false feature, becomes an LLM hallucination, becomes a poisoned training signal, becomes an autonomous agent's wrong action. HSGA's containment strategy assigns a specific holonic mechanism to interrupt each stage of that cascade rather than relying on one general-purpose safeguard at the end.

Cascade Stage	Failure Mode	Holonic Containment Mechanism
Data layer	Sensor noise, calibration drift, source quality degradation	Inbound validation pipeline; data-contract SLAs; ISO 8000 fitness assessment
Algorithmic layer	Inflexible models transform noise into false features	SHACL constraint enforcement; SWRL rule-governed inference
LLM layer	The model finds convincing, erroneous meaning in bad data	Circuit Breaker trip; dimensional MTBH scoring; Intent Delta evaluation

Cascade Stage	Failure Mode	Holonic Containment Mechanism
Recursive learning	Tripped outputs contaminate training data (model collapse)	Sanitization layer excludes every tripped Circuit Breaker output from training
Autonomous agent	The agent executes a physically or strategically wrong action	Simplex Action Gateway; mandatory human authorization for irreversible actions

The trust metric underlying this table is **Mean Time Between Hallucinations (MTBH)** — the quantitative measure of how long the architecture's governance integrity holds between fidelity failures. Critically, MTBH is scored *dimensionally*, across the same five relevancy dimensions used throughout the architecture, rather than as a single aggregate number. A failure in MTBH_WHERE (domain-routing) does not trigger remediation of MTBH_WHAT (concept resolution) in an unrelated domain — each dimension monitors, and can be repaired at, exactly the holonic level it corresponds to. Two companion metrics sharpen this further: the **Value-Drift Coefficient** measures how far an autonomous agent's current objective function has moved from the human intent encoded in its governing ontology, and **Intent Delta** measures the discrepancy between an agent's logical understanding of a task and the physical or operational reality it is actually operating within.

Part VII — Level 4: The Enterprise Holarchy

Federated governance at enterprise scale is, in this architecture, a holarchic organization rather than a management hierarchy. Three tiers each govern at exactly the holonic level that matches their scope of authority, and lower tiers are not subordinate reports to higher ones so much as holons operating at a narrower governance scope within the same protocol.

Governance Tier	Holonic Level Governed	Authority Scope
Enterprise Semantic Governance Council	L4 — Enterprise Holarchy	Strategic priorities, cross-domain conflict resolution, policy approval, investment
Domain Governance Teams	L2 — Domain Holons	Domain ontology development and maintenance, data-quality standards, use-case prioritization
Semantic Governance Center of Excellence	L1 Knowledge + L3 Gateway + Mediation Holon	Platform operations, cross-domain integration, agent development, training and enablement

The same holarchy also explains why the enterprise's decision-making does not simply run “faster” or “slower” at different tiers — it activates different holonic levels and different governance protocols depending on the temporal depth of the decision. An operational adjustment made in seconds to minutes engages only Level 1 and Level 2 pre-validated rules, with no full gateway traversal required. A tactical decision made over hours to days engages Level 2 and Level 3 with full semantic-gateway traversal and cross-domain consistency checks. A strategic decision made over weeks to quarters engages the full Level 4 federated governance council, including LKIF deontic-norm evaluation and cross-domain impact analysis. Each tier gets exactly the governance rigor its decision horizon requires — nothing is either under-governed for its stakes or over-governed for its urgency.

Two supporting mechanisms operate across every tier. Every ISO 8000 data-quality dimension — accuracy, completeness, consistency, timeliness, validity, uniqueness, relevance, traceability — maps directly to a SHACL-enforceable holonic invariant, so a data-quality policy is not a document to consult but a constraint the graph itself enforces; a violation in any one dimension is diagnosable to the exact holonic level and mechanism responsible for it, from inbound rejection to KGCL conflict resolution. And an operations layer — the holarchy's own self-awareness system — gives every one of the four levels, plus the Mediation Holon, a real-time, navigable view of governance state: change events, validation events, trust scores, and lifecycle status are all surfaced as they happen, so governance health is observed from within the holarchy rather than reported on after the fact.

Part VIII — The Self-Evolving Holarchy

A governance architecture that can only enforce what it already knows will fall behind the enterprise it governs. HSGA treats extension and adaptation as first-class, governed capabilities rather than as exceptions to be handled outside the architecture.

Retrieval-augmented generation as a governed extension mechanism

Where SHACL constraints and SWRL rules govern what the holarchy currently knows, retrieval-augmented generation (RAG) is the mechanism by which the holarchy discovers what it *should* know next — new entities, relationships, and ontology classes surfaced by operational experience. A two-level hybrid search implements this: a fast approximate nearest-neighbor search within a domain holon's vector space for pattern recognition, paired with precise symbolic reasoning and SHACL validation to confirm compliance and interpretability. As of Version 3.0, the entity-anchoring step from the Mediation Holon (Part V) runs between these two levels, so retrieved content is bound to a canonical entity before it is ever validated against governance rules. Discovery runs on four governed feedback cycles — new-entity discovery, cross-domain ontology alignment, within-domain ontology enrichment, and anomaly detection and correction — and in every case, what RAG proposes still passes through domain-expert or Domain Governance Team review and a KGCL change record before it becomes part of the governed graph. Discovery never writes directly to canonical state.

SUDA — the cognitive protocol for every AI interaction

The SUDA framework — Sense, Understand, Decide, Act — describes what happens, cognitively, whenever any agent or user interacts with the holarchy: Sense collects and dimensionally tags the interaction (building the initial relevancy context); Understand activates the relevant domain ontology and resolves semantic meaning within its governance context; Decide assembles the Requirement Context Bundle and the Compiled AI Directive, deciding what to retrieve and what to govern; Act passes the governed directive to the AI agent, whose behavior remains bounded by the context assembled in the prior two phases. SUDA and the semantic gateway describe the same traversal from two angles: SUDA is the cognitive model — how the system thinks about context — while the gateway is the governance model — how the system enforces it. Together they define what it means for an AI system to act with holonic comprehension rather than mere data processing.

Observe → Reason → Adapt: closing the design-time / run-time boundary

The INCOSE Knowledge Systems Working Group's Observe-Reason-Adapt loop describes the holarchy's continuous evolution mechanism. Every service call, agent response, MTBH measurement, and KGCL change event is a semantic telemetry signal tagged with full relevancy context (Observe). Reasoning about that telemetry is distributed across every level — SHACL constraints at L1, SWRL rules and domain agents at L2, governance synthesis at L3, enterprise policy review at L4 (Reason). The holarchy then reconfigures itself — ontology updates, change records, revised rules — always constrained by the same governance protocol that prevents adaptation from violating a holonic invariant (Adapt). This closes what has historically been model-based systems engineering's central challenge: the model built

at design time accumulates run-time intelligence and evolves, but always within governance constraints that make the evolution trustworthy. The distinction between modeling the system and operating the system disappears.


Part IX — Governance in Motion

The preceding parts describe governance level by level. This part shows the four governance domains, the automated agent fabric that maintains them, and the human oversight structure that ties every layer together as they actually run — concurrently, continuously, and under consistent human review.

IX.1 Four governance domains, running continuously

Four Governance Domains

Ambient across every layer of the holarchy — enforced at each layer boundary by the Gateway Holon



The four domains run concurrently and continuously, not as sequential phases. All four converge at the Governance Layer and are enforced structurally at every Gateway Holon traversal.

Figure 4: Four Governance Domains — Data, Semantic, Analytics, and AI/Trust govern concurrently, not as sequential phases, and converge structurally at every Gateway Holon traversal.

Data, Semantic, Analytics, and AI/Trust governance are not sequential phases an asset passes through once. Each domain owns its own capability set, its own key performance indicators, and its own tooling, and all four run continuously in parallel — converging at the compiled governance layer described in Part V, and enforced structurally, not procedurally, at every Gateway Holon traversal.

IX.2 The agentic layer: automation that proposes, never commits

A named fabric of automated agent roles maintains every mapping type and governance mechanism described in this paper — from the Lineage Ingestion Agents that keep the Mediation Holon's mapping graph current, to the Rule Harmonization and Rule Projection Agents that compile canonical governance into runtime targets, to a dedicated Trust and Circuit Breaker Agent that scores every model response across the MTBH dimensions.

The Agentic Layer

The agent fabric that maintains every mapping type and executes governance under human supervision



No agent writes directly to the canonical ontology, the mapping graph, or the compiled policy registry — every agent output is a proposal that enters this loop, and HITL review is the only gate through which it becomes a governed, KGCL-recorded change.

Figure 5: The Agentic Layer — the same Observe → Reason → Propose → Human Review → Adapt cycle governs every agent role.


The design constraint that makes this fabric trustworthy is structural, not procedural: no agent writes directly to the canonical ontology, the mapping graph, or the compiled policy registry. Every agent output is a *proposal*, and human-in-the-loop review is the only gate through which a proposal becomes a governed, KGCL-recorded change.

IX.3 One human-in-the-loop spine, not six separate processes

Across every mechanism described in this paper — ontology change, mapping updates, agent proposals, rule harmonization, compiled-policy promotion — the same five-step review cycle applies: Observe, Reason, Propose, Human-in-the-Loop Review, Adapt. These are not six unrelated governance processes bolted onto six different layers. They are one governance spine, applied consistently at every point in the architecture where the system could otherwise learn, update, or retrain silently.

Human-in-the-Loop: One Governance Spine Across the Whole Topology

How HITL guides training and recursive learning across every layer — not a separate process, the same one everywhere



No layer in this architecture learns, updates, or retrains silently. Every candidate change, from an ontology edit to a compiled policy promotion, passes through this same cycle before it becomes governed system state.

Figure 6: Human-in-the-Loop — one governance spine across the whole topology. No layer in this architecture learns, updates, or retrains without passing through it.

The guarantee this buys — no candidate change, from an ontology edit to a compiled policy promotion, becomes governed system state without passing through this same cycle. That consistency is what makes the architecture auditable as a whole rather than auditable layer by layer.

Part X — Standards Foundation

HSGA is not a proprietary format layered on top of open standards; it is built directly on them, and treats every referenced standard as a *holarchic contract* — a protocol established at the enterprise or inter-enterprise level that all holons within scope must conform to. This distinction matters structurally: a standard that imposes schema coupling destroys holonic autonomy, while a standard that defines an interface protocol preserves it. Every standard below is used in the second sense.

Standard	Holonic Role	Applicable Level
OWL 2 (W3C)	Knowledge holon semantic definition language	L1 — Knowledge Holons
SHACL (W3C)	Holon invariant enforcement language; canonical definition-time source of truth	L1/L2; compiled at Mediation Holon
SWRL (W3C)	Holon inference logic — what can be derived	L1/L2; compiled at Mediation Holon
ISO 8000	Knowledge holon quality-invariant specification	L1/L2 — all data holons
LKIF	Deontic norm holon specification language	L1/L2 — normative holons
FIBO (EDM Council)	Financial-domain holon reference ontology	L2 — finance domain holon
STPA/STAMP (MIT)	Safety holonic invariant methodology	L3 — Gateway Holon
SysML v2 / KerML (OMG)	Model-based systems engineering holon design language	L3/L4 — system and enterprise holons
Model Context Protocol (Anthropic)	Holarchic agent communication protocol	L2/L3 — agent and gateway holons
Open Policy Agent / Rego	Compiled policy-decision holonic contract	L3 Gateway / Mediation Holon
dbt + OpenLineage	Bottom-up transformation lineage protocol	Mediation Holon (bottom-up integration)

Part XI — Strategic Value and the Path Forward

The conventional enterprise data paradigm is organized around data management: collecting, storing, processing, and reporting data. HSGA represents a shift to knowledge governance: governing what data means, what it may be used for, and what actions it may authorize — at every level of the enterprise, in every interaction, in real time. That shift is not aspirational; it is structurally achievable through the architecture described in this paper, and it produces four concrete forms of value.

Knowledge as a monetized asset

When semantic assets are governed as knowledge holons with formal identity, quality dimensions, data contracts, and provenance tracking, they acquire the properties that make an asset monetizable: they become discoverable, trustworthy, measurable, and portable across contexts, rather than trapped in one system's schema.

Business processes as intelligent holons

When processes are expressed as governance-grounded process holons — a financial planning process, a budgeting process, a compliance-reporting process — with deontic governance invariants attached, they become able to reason about their own compliance, flag violations, and invoke corrective action without waiting for a human to notice the problem first.

AI-driven automation that operates transparently and safely

When AI agents operate within the governance structure described in this paper — bounded by domain holon invariants, monitored by dimensional trust scoring, protected by Circuit Breaker anti-collapse protocols, grounded against canonical entities by the Mediation Holon, and audited through immutable registry trails — autonomy and trustworthiness stop being in tension. They are the same architectural property, expressed at different levels of the holarchy.

A durable basis for competitive advantage

The organizations that will define the next phase of knowledge-driven enterprise performance are the ones that achieve what Kurt Cagle terms Digital Orchestration — not systems that record what happened, but systems that understand what should happen and can show, structurally, that what happened conformed to it. That is a harder property to copy than a data platform or a model, because it is a property of the governance architecture itself, not of any one tool built on top of it.

Operating principles, in brief

The full architecture is governed by seventeen implementation principles; the six most consequential for leadership evaluating this model are worth stating directly:

- Every semantic asset is a sovereign holon with identity, invariants, and a lifecycle contract before it is trusted with anything — assets without holonic identity are not components of this architecture; they are just data.

- Governance rules are canonical in exactly one place — the graph layer — and compiled everywhere they are enforced. They are never re-authored locally, and every compiled artifact carries a traceable link back to its source rule.
- Trust is monitored dimensionally, not as a single aggregate score, so a failure localizes to the holonic level responsible for it instead of triggering blanket shutdown.
- No automated agent writes directly to governed state. Every agent output is a proposal, and human-in-the-loop review is the only path from proposal to governed change.
- Standards are holarchic contracts between holons, never schemas that force central coupling — the distinction that preserves domain autonomy at enterprise scale.
- The system that governs design must be continuously nourished by the system it governs at run time — this is what closes the model-staleness failure described in Part I for good.

Conclusion

Holonic design, in this architecture, is not a layer added on top of enterprise data management. It is the architectural grammar that makes every governance pattern in this document legible, composable, and trustworthy at enterprise scale — from the smallest governed fact up through the full federated governance ecosystem, with a lateral Mediation Holon holding the whole structure resilient against the ordinary churn of a modern technology stack.

Final principle — *conventional enterprise architectures are built to process data. Holonic architectures are built to govern knowledge. The difference is not one of degree but of kind. An enterprise that builds on holonic foundations builds a system that thinks — not merely one that computes.*

For leadership evaluating this model, the practical proposition is straightforward: HSGA offers a way to scale AI governance without re-centralizing control, a way to make compliance a provable, structural property rather than an asserted one, and a way to let individual domains move at their own pace without fragmenting the enterprise's shared meaning. As AI agents take on a growing share of enterprise decision-making, that is not an incremental improvement to existing data architecture. It is the governance foundation that makes trustworthy AI at enterprise scale possible at all.

Appendix A — Glossary of Holonic Terms

Term	Definition, as used in this architecture
Holon	An entity that is simultaneously a self-contained whole and a composable part of a larger whole.
Holarchy	The nested organizational structure of holons, governed by protocol rather than centralized hierarchy.
Holonic autonomy	The property of a holon to govern itself according to its own invariants without external instruction.
Holonic identity	The persistent self-description of a holon across contexts and lifecycle states (URI/IRI, governance contract, version history).
Holonic invariant	A rule that must be satisfied for a holon to be in a valid governance state (a SHACL constraint or LKIF deontic obligation).
Mediation Holon	The lateral holon, introduced in v3.0, whose scope is the correspondence between representations of knowledge across systems, not the knowledge itself.
Mapping Graph	The Mediation Holon's persistently governed set of correspondence assertions between top-down ontology concepts and bottom-up physical representations.
Compiled Semantic Governance	The pattern separating canonical rule definition (in-graph SHACL/SWRL) from rule execution (runtime-native compiled targets).
Rule Harmonization Engine	The component that detects conflicts, dependencies, and jurisdictional overlaps across the canonical rule graph.
Rule Projection Agent	The component that compiles a canonical SHACL/SWRL invariant into a runtime-native enforcement artifact.
ORSD	Ontology Requirements Specification Document — the design-time governance contract for a knowledge holon.
KGCL	Knowledge Graph Change Language — the run-time, fully attributed change record for a knowledge holon.
MTBH	Mean Time Between Hallucinations — the architecture's dimensional AI-trust health metric.
Value-Drift Coefficient	A measure of how far an autonomous agent's objective function has moved from its original governing intent.
Intent Delta	The measured discrepancy between an AI agent's logical task understanding and its actual operational reality.
Circuit Breaker	The holonic fault-containment protocol that trips and isolates a failure at the level where it occurred.
Anti-collapse mechanism	The sanitization protocol that excludes tripped Circuit Breaker outputs from training data, preventing model collapse.
SUDA	Sense → Understand → Decide → Act — the cognitive protocol for every context-aware holonic interaction.

Term	Definition, as used in this architecture
Digital Orchestration	A system that governs what should happen, not merely one that records what happened (Kurt Cagle).

References

- Wallk, Jeffrey. Holonic Semantic Governance Architecture, v4.2 — Reference Topology & Governance Diagram Edition. The Value Enablement Group, LLC (draft).
- Koestler, Arthur. The Ghost in the Machine (1967). Pan Books — original formulation of holon and holarchy theory.
- Cagle, Kurt. “The World Database: Why Holonic Graphs Are the Missing Layer of Resilience.” LinkedIn Pulse.
- Wallk, Jeffrey. Requirements-Driven Semantic Gateway Architecture (RDSG), v2.0 STPA/STAMP Edition. The Value Enablement Group, LLC.
- Wallk, Jeffrey. Circuit Breaker / MTBH Architecture: Morphism-Grounded Compositional Assurance. The Value Enablement Group, LLC.
- Wallk, Jeffrey. Unified Semantic Governance Framework: Data Quality, Contextual Analytics, and Federated Decision Intelligence, v1.0.
- Wallk, Jeffrey. Semantic Abstraction and Mediation Layer: A Three-Tier Model for Enterprise Knowledge Governance (Design Notes). The Value Enablement Group, LLC.
- Wallk, Jeffrey. Trustworthy AI and Semantic Governance: A Multi-Disciplinary Framework. INCOSE KSWG Presentation.
- Wallk, Jeffrey. INCOSE Knowledge Systems Working Group — Context is King. March 2026 KSWG Presentation.
- OWL 2 Web Ontology Language (W3C Recommendation); SHACL — Shapes Constraint Language (W3C); SWRL — Semantic Web Rule Language (W3C).
- ISO 8000 Data Quality Standards; ISO/IEC/IEEE 24765 Systems and Software Engineering Vocabulary.
- LKIF — Legal Knowledge Interchange Format (Estrella Project); FIBO — Financial Industry Business Ontology (EDM Council).
- STPA/STAMP — Systems Theoretic Process Analysis (MIT STAMP Project, Nancy Leveson); OMG SysML v2 / KerML Specification.
- Model Context Protocol (MCP) — Anthropic; Open Policy Agent (OPA) / Rego Policy Language; OpenLineage — Open Standard for Data Lineage; dbt (data build tool).

— End of Document —

The Value Enablement Group, LLC | Jeffrey Wallk